

HW2-R

Nima Kelidari

2023-05-02

I decided to use the second dataset which data is 50-50 for diabetic and non-diabetic persons. this balance will help us to decrease bias in our models. And it is cleaner and has less same row and NA rows so much. (But we try to detect and delete them anyway)

```
#install.packages('caret')
#install.packages('pROC')
#install.packages("ggplot2")
#install.packages("tidyverse")
#install.packages("leaps")
#install.packages("MASS")
#install.packages("tree")
library("tree")
```

```
## Warning: package 'tree' was built under R version 4.2.3
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.2.3
```

```
## Warning: package 'ggplot2' was built under R version 4.2.3
```

```
## Warning: package 'tibble' was built under R version 4.2.3
```

```
## Warning: package 'tidyr' was built under R version 4.2.3
```

```
## Warning: package 'readr' was built under R version 4.2.3
```

```
## Warning: package 'purrr' was built under R version 4.2.3
```

```
## Warning: package 'dplyr' was built under R version 4.2.3
```

```
## Warning: package 'stringr' was built under R version 4.2.3
```

```
## Warning: package 'forcats' was built under R version 4.2.3
```

```
## Warning: package 'lubridate' was built under R version 4.2.3
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.2      v readr      2.1.4
```

```
## v forcats   1.0.0      v stringr   1.5.0
```

```
## v ggplot2   3.4.2      v tibble    3.2.1
```

```
## v lubridate 1.9.2      v tidyr     1.3.0
```

```
## v purrr     1.0.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```

library(leaps)

## Warning: package 'leaps' was built under R version 4.2.3
library(MASS)

##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##      select
library(ggplot2)
library(pROC)

## Warning: package 'pROC' was built under R version 4.2.3
## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##      cov, smooth, var
library(caret)

## Warning: package 'caret' was built under R version 4.2.3
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##      lift
diabets = read.csv("D:/Terme8/Regression/HW2/Data/d2.csv")

#diabets

diabets <- diabets[!duplicated(diabets), ]
diabets <- na.omit(diabets)
diabets$output <- as.factor(diabets$Diabetes_binary)

#diabets

```

PART 1


Yes, by using the proper features and model, we surely can get a nice prediction and then in the next part, use it for inference. The main reason for this is the data given to us, has some features which probably have some relationship with the probability of whether one has diabetes or not. We can show it by plotting those features by the diabetes factor. For example, we plotted BMI, high blood pressure, high cholesterol, general health, age, education, and income by diabetes factor in the type of histogram density, histogram, and box plot. Here we can see the different distribution of these factors for two groups of diabetic and non-diabetic persons. Moreover, we fit a logistic regression to show t-value is so big for this problem and that these factors and responses are correlated strongly. At last, we fit a model by these factors and see almost all of them, have

strongly related to the diabetes factor. So at least some of them can be used for the prediction of whether one has diabetes or not. In the second part, we introduce some of these features as good ones for prediction.

```
model <- glm(Diabetes_binary ~ BMI, data = diabetics)
summary(model)
```

```
##
## Call:
## glm(formula = Diabetes_binary ~ BMI, data = diabetics)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.86763  -0.44918   0.05137   0.47091   0.83051
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.0902293  0.0078552  -11.49  <2e-16 ***
## BMI          0.0199781  0.0002551   78.33  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.2295464)
##
##      Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 15851  on 69055  degrees of freedom
## AIC: 94352
##
## Number of Fisher Scoring iterations: 2
```


```
Diabetstic = as.factor(diabetics$Diabetes_binary)
# Density areas without lines
ggplot(diabetics, aes(x = BMI, fill = Diabetstic, colours = Diabetstic)) +
  geom_density(alpha = 0.4, color = NA)
```



```
model <- glm(Diabetes_binary ~ HighBP, data = diabetics)
summary(model)
```

```
##
## Call:
## glm(formula = Diabetes_binary ~ HighBP, data = diabetics)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.6694  -0.2935   0.3306   0.3306   0.7065
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.293549   0.002697  108.9   <2e-16 ***
## HighBP       0.375830   0.003568  105.3   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.215343)
##
##      Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 14871  on 69055  degrees of freedom
## AIC: 89941
##
## Number of Fisher Scoring iterations: 2
```

```
Diabetstic = as.factor(diabets$Diabetes_binary)
# Density areas without lines
ggplot(diabets, aes(x = HighBP, fill = Diabetstic, colours = Diabetstic)) +
  geom_histogram(alpha = 0.4, position = "identity", bins = 3)
```




```
(legend.position = "none")
```

```
model <- glm(Diabetes_binary ~ HighChol, data = diabets)
summary(model)
```

```
##
## Call:
## glm(formula = Diabetes_binary ~ HighChol, data = diabets)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.6404  -0.3584   0.3596   0.3596   0.6416
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.358443    0.002667  134.42  <2e-16 ***
## HighChol     0.281915    0.003658   77.06  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.2301479)
##
```

```
## Null deviance: 17260 on 69056 degrees of freedom
## Residual deviance: 15893 on 69055 degrees of freedom
## AIC: 94532
##
## Number of Fisher Scoring iterations: 2
```

```
Diabetstic = as.factor(diabets$Diabetes_binary)
# Density areas without lines
ggplot(diabets, aes(x = HighChol, fill = Diabetstic, colours = Diabetstic)) +
  geom_histogram(alpha = 0.4, position = "identity", bins = 3)
```




```
model <- glm(Diabetes_binary ~ GenHlth, data = diabets)
summary(model)
```

```
##
## Call:
## glm(formula = Diabetes_binary ~ GenHlth, data = diabets)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.8905  -0.3537   0.1095   0.4674   0.8253
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.004206   0.004840  -0.869   0.385
## GenHlth      0.178943   0.001576 113.520 <2e-16 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.2106319)
##
##      Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 14545  on 69055  degrees of freedom
## AIC: 88413
##
## Number of Fisher Scoring iterations: 2
```

```
Diabetstic = as.factor(diabets$Diabetes_binary)
# Density areas without lines
ggplot(diabets, aes(x = GenHlth, fill = Diabetstic , colours = Diabetstic)) +
  geom_histogram(alpha = 1, position = "dodge",bins = 9)
```




```
model <- glm(Diabetes_binary ~ Age, data = diabets)
summary(model)
```

```
##
## Call:
## glm(formula = Diabetes_binary ~ Age, data = diabets)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.7193  -0.4792   0.2807   0.4247   0.8569
##
## Coefficients:
```

```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.095058   0.005803   16.38  <2e-16 ***
## Age         0.048021   0.000640   75.03  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.2310996)
##
## Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 15959  on 69055  degrees of freedom
## AIC: 94817
##
## Number of Fisher Scoring iterations: 2
```

```
Diabetstic = as.factor(diabets$Diabetes_binary)
# Density areas without lines
ggplot(diabets, aes(x = Age, fill = Diabetstic , colours = Diabetstic)) +
  geom_boxplot(alpha = 1)
```




```
model <- glm(Diabetes_binary ~ Education, data = diabets)
summary(model)
```

```
##
## Call:
## glm(formula = Diabetes_binary ~ Education, data = diabets)
##
## Deviance Residuals:
```



```
##      Min      1Q   Median      3Q      Max
## -0.8085 -0.5006  0.2685   0.4994  0.5764
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.885516   0.009138  96.91  <2e-16 ***
## Education   -0.076992   0.001825 -42.19  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.2436587)
##
## Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 16826  on 69055  degrees of freedom
## AIC: 98472
##
## Number of Fisher Scoring iterations: 2
```

```
Diabetstic = as.factor(diabets$Diabetes_binary)
# Density areas without lines
ggplot(diabets, aes(x = Education, fill = Diabetstic , colours = Diabetstic)) +
  geom_histogram(alpha = 0.4, position = "identity", bins = 8)
```




```
model <- glm(Diabetes_binary ~ Income, data = diabets)
summary(model)
```

```
##
```

```
## Call:
## glm(formula = Diabetes_binary ~ Income, data = diabetics)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.7357  -0.4423   0.2643   0.5088   0.6066
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.7846393  0.0051740  151.65  <2e-16 ***
## Income      -0.0489101  0.0008544  -57.24  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.2386164)
##
##      Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 16478  on 69055  degrees of freedom
## AIC: 97028
##
## Number of Fisher Scoring iterations: 2
```

```
Diabetstic = as.factor(diabetics$Diabetes_binary)
# Density areas without lines
ggplot(diabetics, aes(x = Income, fill = Diabetstic, colours = Diabetstic)) +
  geom_histogram(alpha = 1, position = "dodge", bins = 6)
```



```

model <- glm(Diabetes_binary ~ Income+Education+Age+GenHlth+HighChol+HighBP+BMI, data = diabetes)
summary(model)

##
## Call:
## glm(formula = Diabetes_binary ~ Income + Education + Age + GenHlth +
##      HighChol + HighBP + BMI, data = diabetes)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.62192  -0.34359   0.03701   0.34423   1.16969
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.4658267  0.0134692  -34.59  < 2e-16 ***
## Income      -0.0091642  0.0008643  -10.60  < 2e-16 ***
## Education   -0.0077290  0.0017604   -4.39  1.13e-05 ***
## Age         0.0280524  0.0006106   45.94  < 2e-16 ***
## GenHlth     0.1050222  0.0016701   62.88  < 2e-16 ***
## HighChol    0.1139753  0.0034549   32.99  < 2e-16 ***
## HighBP     0.1633319  0.0037298   43.79  < 2e-16 ***
## BMI         0.0122764  0.0002384   51.51  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1775591)
##
##      Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 12260  on 69049  degrees of freedom
## AIC: 76624
##
## Number of Fisher Scoring iterations: 2
Diabetstic = as.factor(diabetes$Diabetes_binary)

```

PART 2

We can answer this question by feat a model like logistic regression on whole data and all of the features, then detect those features which have big enough t-values. Here we can see BMI, high blood pressure, high cholesterol, general health, age, cholesterol check, history of heart attack, sex, age, heavy alcohol consumption, and income have an absolute value of t-value more than 10; We can detect these factors as important factors compared to others and check them more carefully. we will see how we can select features more reliably by feature selection in part 3.

```

md <- glm(Diabetes_binary ~., data = diabetes)
summary(md)

##
## Call:
## glm(formula = Diabetes_binary ~ ., data = diabetes)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.554e-14  -4.219e-15  -4.892e-16   4.161e-15   1.920e-14
##
## Coefficients:

```

```
##               Estimate Std. Error   t value Pr(>|t|)
## (Intercept)      7.367e-15  2.229e-16  3.305e+01 < 2e-16 ***
## HighBP          -1.949e-15  4.689e-17 -4.157e+01 < 2e-16 ***
## HighChol        -1.342e-15  4.323e-17 -3.104e+01 < 2e-16 ***
## CholCheck       -2.073e-15  1.287e-16 -1.611e+01 < 2e-16 ***
## BMI             -1.464e-16  3.050e-18 -4.799e+01 < 2e-16 ***
## Smoker           4.170e-17  4.110e-17  1.015e+00 0.310308
## Stroke          -3.051e-16  8.477e-17 -3.598e+00 0.000320 ***
## HeartDiseaseorAttack -6.008e-16  6.042e-17 -9.944e+00 < 2e-16 ***
## PhysActivity     7.286e-17  4.634e-17  1.572e+00 0.115872
## Fruits           4.770e-17  4.241e-17  1.125e+00 0.260690
## Veggies          1.339e-16  5.043e-17  2.654e+00 0.007945 **
## HvyAlcoholConsump 1.548e-15  9.834e-17  1.574e+01 < 2e-16 ***
## AnyHealthcare    -1.709e-16  9.879e-17 -1.730e+00 0.083687 .
## NoDocbcCost      1.806e-17  7.173e-17  2.520e-01 0.801257
## GenHlth         -1.294e-15  2.483e-17 -5.212e+01 < 2e-16 ***
## MentHlth         1.096e-17  2.702e-18  4.055e+00 5.02e-05 ***
## PhysHlth         1.940e-17  2.540e-18  7.639e+00 2.22e-14 ***
## DiffWalk        -3.559e-16  5.687e-17 -6.257e+00 3.95e-10 ***
## Sex              -5.652e-16  4.153e-17 -1.361e+01 < 2e-16 ***
## Age              -3.069e-16  8.088e-18 -3.794e+01 < 2e-16 ***
## Education        7.614e-17  2.203e-17  3.456e+00 0.000548 ***
## Income           1.340e-16  1.122e-17  1.194e+01 < 2e-16 ***
## output1          1.000e+00  4.733e-17  2.113e+16 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 2.706099e-29)
##
##      Null deviance: 1.7260e+04  on 69056  degrees of freedom
## Residual deviance: 1.8681e-24  on 69034  degrees of freedom
## AIC: -4346532
##
## Number of Fisher Scoring iterations: 1
```

PART 3

Yes, absolutely we can do it by feature selection methods. As best subset selection can be very time-consuming and computationally heavy in practice (in real cases), we just use two methods forward and backward selection by AIC. In each one, we want to use 5 features maximum utmost. so we make a full model and null model first and try to extract the five most important features. By each of these two methods, we can see five features below BMI, high blood pressure, high cholesterol, general health, and age, have the most importance in prediction as we guss before in part 2; in other words, these have most dependency to pred factor. On another hand, these do have not a significant relation to each other, because, in a model that we created before, they didn't decrease each other t-value strongly. so we use these five features for use in the next parts.

```
full.model <- glm(Diabetes_binary ~.-output, data = diabet)
step.model <- stepAIC(full.model, direction = "both", trace = FALSE)
summary(step.model)
```

```
##
## Call:
## glm(formula = Diabetes_binary ~ HighBP + HighChol + CholCheck +
##      BMI + Stroke + HeartDiseaseorAttack + PhysActivity + Veggies +
##      HvyAlcoholConsump + AnyHealthcare + GenHlth + MentHlth +
##      PhysHlth + DiffWalk + Sex + Age + Education + Income, data = diabet)
```

```
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.55229  -0.33889   0.03907   0.33968   1.25189
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -0.5923088   0.0173947  -34.051 < 2e-16 ***
## HighBP          0.1540678   0.0037242   41.369 < 2e-16 ***
## HighChol        0.1088292   0.0034467   31.575 < 2e-16 ***
## CholCheck       0.1667005   0.0103164   16.159 < 2e-16 ***
## BMI             0.0118230   0.0002408   49.091 < 2e-16 ***
## Stroke          0.0252910   0.0068154    3.711 0.000207 ***
## HeartDiseaseorAttack 0.0482946   0.0048483    9.961 < 2e-16 ***
## PhysActivity    -0.0063787   0.0037116   -1.719 0.085698 .
## Veggies         -0.0118939   0.0039675   -2.998 0.002720 **
## HvyAlcoholConsump -0.1252300   0.0078558  -15.941 < 2e-16 ***
## AnyHealthcare    0.0139394   0.0078103    1.785 0.074306 .
## GenHlth         0.1037024   0.0019532   53.093 < 2e-16 ***
## MentHlth        -0.0008754   0.0002161   -4.051 5.11e-05 ***
## PhysHlth        -0.0015695   0.0002039   -7.697 1.41e-14 ***
## DiffWalk        0.0286401   0.0045682    6.269 3.65e-10 ***
## Sex             0.0455012   0.0032980   13.797 < 2e-16 ***
## Age             0.0247627   0.0006350   38.996 < 2e-16 ***
## Education       -0.0060606   0.0017639   -3.436 0.000591 ***
## Income          -0.0107495   0.0008965  -11.990 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1749733)
##
##      Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 12080  on 69038  degrees of freedom
## AIC: 75622
##
## Number of Fisher Scoring iterations: 2

train.control <- trainControl(method = "cv", number = 10)
step.model <- train(Diabetes_binary ~.-output, data = diabet,
                    method = "leapBackward",
                    tuneGrid = data.frame(nvmax = 1:5),
                    trControl = train.control
                    )

## Warning in train.default(x, y, weights = w, ...): You are trying to do
## regression and your outcome only has two possible values Are you trying to do
## classification? If so, use a 2 level factor as your outcome column.

step.model$results
```

	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	0.4589429	0.1573782	0.4212581	0.002282365	0.008347025	0.001917050
## 2	2	0.4451980	0.2072301	0.3970476	0.002754008	0.009545242	0.003708062
## 3	3	0.4338239	0.2471322	0.3818945	0.003863398	0.013234345	0.002302323
## 4	4	0.4252467	0.2765826	0.3678374	0.003860724	0.013067806	0.003697475

```
## 5      5 0.4219822 0.2876314 0.3631322 0.003694688 0.012444030 0.003502273
```

```
step.model$bestTune
```

```
## nvmax
```

```
## 5      5
```

```
summary(step.model$finalModel)
```

```
## Subset selection object
```

```
## 21 Variables (and intercept)
```

```
##              Forced in Forced out
```

```
## HighBP              FALSE      FALSE
```

```
## HighChol            FALSE      FALSE
```

```
## CholCheck           FALSE      FALSE
```

```
## BMI                 FALSE      FALSE
```

```
## Smoker              FALSE      FALSE
```

```
## Stroke              FALSE      FALSE
```

```
## HeartDiseaseorAttack FALSE      FALSE
```

```
## PhysActivity        FALSE      FALSE
```

```
## Fruits              FALSE      FALSE
```

```
## Veggies             FALSE      FALSE
```

```
## HvyAlcoholConsump   FALSE      FALSE
```

```
## AnyHealthcare       FALSE      FALSE
```

```
## NoDocbcCost         FALSE      FALSE
```

```
## GenHlth             FALSE      FALSE
```

```
## MentHlth            FALSE      FALSE
```

```
## PhysHlth            FALSE      FALSE
```

```
## DiffWalk            FALSE      FALSE
```

```
## Sex                 FALSE      FALSE
```

```
## Age                 FALSE      FALSE
```

```
## Education           FALSE      FALSE
```

```
## Income              FALSE      FALSE
```

```
## 1 subsets of each size up to 5
```

```
## Selection Algorithm: backward
```

```
##              HighBP HighChol CholCheck BMI Smoker Stroke HeartDiseaseorAttack
```

```
## 1 ( 1 ) " " " " " " " " " " " "
```

```
## 2 ( 1 ) " " " " " " " " " " " "
```

```
## 3 ( 1 ) " " " " " " "*" " " " " "
```

```
## 4 ( 1 ) "*" " " " " "*" " " " " "
```

```
## 5 ( 1 ) "*" "*" " " "*" " " " " "
```

```
##              PhysActivity Fruits Veggies HvyAlcoholConsump AnyHealthcare
```

```
## 1 ( 1 ) " " " " " " " " " " "
```

```
## 2 ( 1 ) " " " " " " " " " " "
```

```
## 3 ( 1 ) " " " " " " " " " " "
```

```
## 4 ( 1 ) " " " " " " " " " " "
```

```
## 5 ( 1 ) " " " " " " " " " " "
```

```
##              NoDocbcCost GenHlth MentHlth PhysHlth DiffWalk Sex Age Education
```

```
## 1 ( 1 ) " " "*" " " " " " " " " " "
```

```
## 2 ( 1 ) " " "*" " " " " " " " " "*" " "
```

```
## 3 ( 1 ) " " "*" " " " " " " " " "*" " "
```

```
## 4 ( 1 ) " " "*" " " " " " " " " "*" " "
```

```
## 5 ( 1 ) " " "*" " " " " " " " " "*" " "
```

```
##              Income
```

```
## 1 ( 1 ) " "
```

```
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "
```

```
coef(step.model$finalModel, 5)
```

```
## (Intercept)      HighBP      HighChol      BMI      GenHlth      Age
## -0.58663321  0.16583644  0.11341426  0.01235041  0.11299873  0.02862784
```

```
mod <- glm(Diabetes_binary ~ HighBP + HighChol + BMI + GenHlth + Age,
  data = diabetes)
summary(mod)
```

```
##
## Call:
## glm(formula = Diabetes_binary ~ HighBP + HighChol + BMI + GenHlth +
##      Age, data = diabetes)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.60267  -0.34499   0.04069   0.34305   1.16937
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.5866332  0.0090232  -65.01  <2e-16 ***
## HighBP      0.1658364  0.0037307   44.45  <2e-16 ***
## HighChol    0.1134143  0.0034596   32.78  <2e-16 ***
## BMI         0.0123504  0.0002386   51.76  <2e-16 ***
## GenHlth     0.1129987  0.0015729   71.84  <2e-16 ***
## Age         0.0286278  0.0006100   46.93  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1780615)
##
##      Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 12295  on 69051  degrees of freedom
## AIC: 76817
##
## Number of Fisher Scoring iterations: 2
```

```
full.model <- glm(Diabetes_binary ~.-output, data = diabetes)
step.model <- stepAIC(full.model, direction = "both", trace = FALSE)
summary(step.model)
```

```
##
## Call:
## glm(formula = Diabetes_binary ~ HighBP + HighChol + CholCheck +
##      BMI + Stroke + HeartDiseaseorAttack + PhysActivity + Veggies +
##      HvyAlcoholConsump + AnyHealthcare + GenHlth + MentHlth +
##      PhysHlth + DiffWalk + Sex + Age + Education + Income, data = diabetes)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.55229  -0.33889   0.03907   0.33968   1.25189
```

```
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -0.5923088  0.0173947 -34.051 < 2e-16 ***
## HighBP         0.1540678  0.0037242  41.369 < 2e-16 ***
## HighChol       0.1088292  0.0034467  31.575 < 2e-16 ***
## CholCheck      0.1667005  0.0103164  16.159 < 2e-16 ***
## BMI            0.0118230  0.0002408  49.091 < 2e-16 ***
## Stroke         0.0252910  0.0068154   3.711 0.000207 ***
## HeartDiseaseorAttack 0.0482946  0.0048483   9.961 < 2e-16 ***
## PhysActivity   -0.0063787  0.0037116  -1.719 0.085698 .
## Veggies        -0.0118939  0.0039675  -2.998 0.002720 **
## HvyAlcoholConsump -0.1252300  0.0078558 -15.941 < 2e-16 ***
## AnyHealthcare   0.0139394  0.0078103   1.785 0.074306 .
## GenHlth        0.1037024  0.0019532  53.093 < 2e-16 ***
## MentHlth       -0.0008754  0.0002161  -4.051 5.11e-05 ***
## PhysHlth       -0.0015695  0.0002039  -7.697 1.41e-14 ***
## DiffWalk       0.0286401  0.0045682   6.269 3.65e-10 ***
## Sex            0.0455012  0.0032980  13.797 < 2e-16 ***
## Age            0.0247627  0.0006350  38.996 < 2e-16 ***
## Education      -0.0060606  0.0017639  -3.436 0.000591 ***
## Income         -0.0107495  0.0008965 -11.990 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1749733)
##
##    Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 12080  on 69038  degrees of freedom
## AIC: 75622
##
## Number of Fisher Scoring iterations: 2

set.seed(123)
train.control <- trainControl(method = "cv", number = 10)
step.model <- train(Diabetes_binary ~.-output, data = diabet,
                    method = "leapForward",
                    tuneGrid = data.frame(nvmax = 1:5),
                    trControl = train.control
                    )

## Warning in train.default(x, y, weights = w, ...): You are trying to do
## regression and your outcome only has two possible values Are you trying to do
## classification? If so, use a 2 level factor as your outcome column.

step.model$results

##   nvmax      RMSE Rsquared      MAE      RMSESD RsquaredSD      MAESD
## 1     1 0.4589449 0.1573952 0.4212585 0.002299933 0.008421811 0.001758425
## 2     2 0.4397916 0.2262143 0.3868186 0.002309520 0.008198404 0.001832603
## 3     3 0.4334785 0.2482743 0.3780406 0.002330834 0.008151682 0.001870903
## 4     4 0.4252604 0.2765075 0.3678396 0.002230627 0.007686012 0.002104354
## 5     5 0.4219997 0.2875429 0.3631400 0.002065983 0.007077088 0.002079634

step.model$bestTune
```



```
## nvmax
## 5 5
```

```
summary(step.model$finalModel)
```

```
## Subset selection object
```

```
## 21 Variables (and intercept)
```

	Forced in	Forced out
## HighBP	FALSE	FALSE
## HighChol	FALSE	FALSE
## CholCheck	FALSE	FALSE
## BMI	FALSE	FALSE
## Smoker	FALSE	FALSE
## Stroke	FALSE	FALSE
## HeartDiseaseorAttack	FALSE	FALSE
## PhysActivity	FALSE	FALSE
## Fruits	FALSE	FALSE
## Veggies	FALSE	FALSE
## HvyAlcoholConsump	FALSE	FALSE
## AnyHealthcare	FALSE	FALSE
## NoDocbcCost	FALSE	FALSE
## GenHlth	FALSE	FALSE
## MentHlth	FALSE	FALSE
## PhysHlth	FALSE	FALSE
## DiffWalk	FALSE	FALSE
## Sex	FALSE	FALSE
## Age	FALSE	FALSE
## Education	FALSE	FALSE
## Income	FALSE	FALSE

```
## 1 subsets of each size up to 5
```

```
## Selection Algorithm: forward
```

	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorAttack
## 1 (1)	" "	" "	" "	" "	" "	" "	" "
## 2 (1)	"*	" "	" "	" "	" "	" "	" "
## 3 (1)	"*	" "	" "	"*	" "	" "	" "
## 4 (1)	"*	" "	" "	"*	" "	" "	" "
## 5 (1)	"*	"*	" "	"*	" "	" "	" "

	PhysActivity	Fruits	Veggies	HvyAlcoholConsump	AnyHealthcare
## 1 (1)	" "	" "	" "	" "	" "
## 2 (1)	" "	" "	" "	" "	" "
## 3 (1)	" "	" "	" "	" "	" "
## 4 (1)	" "	" "	" "	" "	" "
## 5 (1)	" "	" "	" "	" "	" "

	NoDocbcCost	GenHlth	MentHlth	PhysHlth	DiffWalk	Sex	Age	Education
## 1 (1)	" "	"*	" "	" "	" "	" "	" "	" "
## 2 (1)	" "	"*	" "	" "	" "	" "	" "	" "
## 3 (1)	" "	"*	" "	" "	" "	" "	" "	" "
## 4 (1)	" "	"*	" "	" "	" "	" "	"*	" "
## 5 (1)	" "	"*	" "	" "	" "	" "	"*	" "

	Income
## 1 (1)	" "
## 2 (1)	" "
## 3 (1)	" "
## 4 (1)	" "
## 5 (1)	" "

```

coef(step.model$finalModel, 5)

## (Intercept)      HighBP      HighChol      BMI      GenHlth      Age
## -0.58663321  0.16583644  0.11341426  0.01235041  0.11299873  0.02862784

mod <- glm(Diabetes_binary ~ HighBP + HighChol + BMI + GenHlth + Age,
  data = diabets)
summary(mod)

##
## Call:
## glm(formula = Diabetes_binary ~ HighBP + HighChol + BMI + GenHlth +
##      Age, data = diabets)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.60267  -0.34499   0.04069   0.34305   1.16937
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.5866332  0.0090232  -65.01  <2e-16 ***
## HighBP      0.1658364  0.0037307   44.45  <2e-16 ***
## HighChol    0.1134143  0.0034596   32.78  <2e-16 ***
## BMI         0.0123504  0.0002386   51.76  <2e-16 ***
## GenHlth     0.1129987  0.0015729   71.84  <2e-16 ***
## Age         0.0286278  0.0006100   46.93  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1780615)
##
##      Null deviance: 17260  on 69056  degrees of freedom
## Residual deviance: 12295  on 69051  degrees of freedom
## AIC: 76817
##
## Number of Fisher Scoring iterations: 2

```

PART 4

We will use the five features that are mentioned in the last part. If we use all of the features or more than five, our efficiency for the trained model will drop. In this part, we use 70 percent of the data for the training and 30 of the data for the test. Here we will not use a validation set and we prefer to use cross-validation, more specifically 5fold-cv for it, it's time-consuming. Then we will try to train the model, tune its hyperparameters, and then test it and get accuracy and precision, and recall from them, then make the confusion matrix for each model based on the test set. We here will use QDA, LDA, Randomforest, Neural Network, Tree, Linear SVM, Logistic regression, and KNN models. After the test, we will save the accuracy of each model and show and compare them at last and choose one of them. (Here we can see the Tree model had excellent performance besides its simplicity)

```

diabets$Diabetes_binary <- as.factor(diabets$Diabetes_binary)
preProcess <- c("center","scale")
i <- createDataPartition(y = diabets$Diabetes_binary, times = 1, p = 0.7, list = FALSE)
training_set <- diabets[i,]
test_set <- diabets[-i,]
trControl <- trainControl(method = "cv", number = 5)

```

```
model <- train(Diabetes_binary ~ Age+GenHlth+HighChol+HighBP+BMI, method='qda', data = training_set, me
summary(model)
```

```
##           Length Class      Mode
## prior           2    -none-   numeric
## counts          2    -none-   numeric
## means          10    -none-   numeric
## scaling         50    -none-   numeric
## ldet            2    -none-   numeric
## lev             2    -none-   character
## N               1    -none-   numeric
## call            3    -none-   call
## xNames          5    -none-   character
## problemType     1    -none-   character
## tuneValue       1    data.frame list
## obsLevels       2    -none-   character
## param           0    -none-   list
```

```
test_set$pred <- predict(model, test_set)
test_set$factor_pred <- as.factor(test_set$pred)
test_set$factor_truth <- as.factor(test_set$Diabetes_binary)
precision <- posPredValue(test_set$factor_truth, test_set$factor_pred)
recall <- sensitivity(test_set$factor_truth, test_set$factor_pred)
cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table
test_set$pred <- predict(model, test_set, probability=TRUE)
#-----
print(paste("Precision: ", precision))
```

```
## [1] "Precision: 0.688751472320377"
```

```
print(paste("Accuracy: ", accuracy))
```

```
## [1] "Accuracy: 0.733793502920307"
```

```
print(paste("Recall: ", recall))
```

```
## [1] "Recall: 0.749599401773315"
```

```
acc_qda <- accuracy
```

```
#-----
```

```
Reference <- factor(c(0, 1, 0, 1))
```


```
Prediction <- factor(c(0, 0, 1, 1))
```

```
Y <- array(cmat)
```

```
df <- data.frame(Prediction, Reference, Y)
```

```
library(ggplot2)
```

```
ggplot(data = df, mapping = aes(x = Reference, y = Prediction)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%1.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow")
```



```
model <- train(Diabetes_binary ~ Age+GenHlth+HighChol+HighBP+BMI, method='rf', data = training_set, method = 'rf')
summary(model)
```

```
##           Length Class      Mode
## call           4 -none-    call
## type           1 -none- character
## predicted     48340 factor   numeric
## err.rate       1500 -none-    numeric
## confusion        6 -none-    numeric
## votes         96680 matrix   numeric
## oob.times      48340 -none-    numeric
## classes        2 -none- character
## importance       5 -none-    numeric
## importanceSD     0 -none-    NULL
## localImportance  0 -none-    NULL
## proximity       0 -none-    NULL
## ntree           1 -none-    numeric
## mtry            1 -none-    numeric
## forest         14 -none-    list
## y              48340 factor   numeric
## test           0 -none-    NULL
## inbag           0 -none-    NULL
## xNames          5 -none- character
## problemType     1 -none- character
## tuneValue       1 data.frame list
## obsLevels       2 -none- character
```

```

## param          0 -none-      list
test_set$pred <- predict(model, test_set)
test_set$factor_pred <- as.factor(test_set$pred)
test_set$factor_truth <- as.factor(test_set$Diabetes_binary)
precision <- posPredValue(test_set$factor_truth, test_set$factor_pred)
recall <- sensitivity(test_set$factor_truth, test_set$factor_pred)
cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table
test_set$pred <- predict(model, test_set, probability=TRUE)
#-----
print(paste("Precision: ", precision))

## [1] "Precision:  0.692775814683942"

print(paste("Accuracy: ", accuracy))

## [1] "Accuracy:  0.739441038760438"


print(paste("Recall: ", recall))

## [1] "Recall:  0.756808921295303"

acc_rf <- accuracy
#-----
Reference <- factor(c(0, 1, 0, 1))
Prediction <- factor(c(0, 0, 1, 1))
Y <- array(cmat)
df <- data.frame(Prediction, Reference, Y)

library(ggplot2)
ggplot(data = df, mapping = aes(x = Reference, y = Prediction)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%1.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow")

```



1

```
model <- train(Diabetes_binary ~ Age+GenHlth+HighChol+HighBP+BMI, method='lda', data = training_set, me
summary(model)
```

```
##           Length Class      Mode
## prior          2    -none-  numeric
## counts          2    -none-  numeric
## means         10    -none-  numeric
## scaling         5    -none-  numeric
## lev            2    -none-  character
## svd             1    -none-  numeric
## N              1    -none-  numeric
## call           4    -none-  call
## xNames          5    -none-  character
## problemType     1    -none-  character
## tuneValue        1  data.frame list
## obsLevels        2    -none-  character
## param           1    -none-  list
```

```
test_set$pred <- predict(model, test_set)
test_set$factor_pred <- as.factor(test_set$pred)
test_set$factor_truth <- as.factor(test_set$Diabetes_binary)
precision <- posPredValue(test_set$factor_truth, test_set$factor_pred)
recall <- sensitivity(test_set$factor_truth, test_set$factor_pred)
cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table
```

```

test_set$pred <- predict(model, test_set,probability=TRUE)
#-----
print(paste("Precision: ",precision))

## [1] "Precision:  0.703572830781311"

print(paste("Accuracy: ",accuracy))

## [1] "Accuracy:  0.736496596997635"


print(paste("Recall: ",recall))

## [1] "Recall:  0.746122618923701"

acc_lda <- accuracy
#-----
Reference <- factor(c(0, 1, 0, 1))
Prediction  <- factor(c(0, 0, 1, 1))
Y <- array(cmat)
df <- data.frame(Prediction, Reference, Y)

library(ggplot2)
ggplot(data = df, mapping = aes(x = Reference, y = Prediction)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%1.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow")

```



```
model <- train(Diabetes_binary ~ Age+GenHlth+HighChol+HighBP+BMI, method='nnet', data = training_set, m
```

```
## # weights: 8
## initial value 27475.631598
## iter 10 value 21038.517775
## iter 20 value 20235.862096
## iter 30 value 20198.696705
## iter 40 value 20184.071335
## iter 50 value 20182.863189
## final value 20182.654174
## converged
## # weights: 22
## initial value 27256.022350
## iter 10 value 20413.645882
## iter 20 value 20240.068026
## iter 30 value 20068.923616
## iter 40 value 20036.505684
## iter 50 value 20026.392762
## iter 60 value 20023.193424
## iter 70 value 20019.805360
## iter 80 value 20018.188481
## iter 90 value 20017.270463
## final value 20017.251381
## converged
## # weights: 36
## initial value 29722.763310
## iter 10 value 20462.328597
## iter 20 value 20256.256665
## iter 30 value 20165.295735
## iter 40 value 20079.760959
## iter 50 value 20028.191547
## iter 60 value 20007.725222
## iter 70 value 19992.919793
## iter 80 value 19989.382511
## iter 90 value 19984.559570
## iter 100 value 19981.601046
## final value 19981.601046
## stopped after 100 iterations
## # weights: 8
## initial value 27373.629310
## iter 10 value 20710.032272
## iter 20 value 20265.636178
## iter 30 value 20202.210224
## iter 40 value 20197.636791
## iter 50 value 20197.587775
## final value 20197.586990
## converged
## # weights: 22
## initial value 27499.391189
## iter 10 value 20276.846938
## iter 20 value 20170.312025
## iter 30 value 20087.787201
## iter 40 value 20047.325060
## iter 50 value 20042.308233
```



```

## iter 60 value 20033.504110
## iter 70 value 20014.760232
## iter 80 value 20011.726617
## iter 90 value 20010.405556
## iter 100 value 20010.352349
## final value 20010.352349
## stopped after 100 iterations
## # weights: 36
## initial value 28724.636485
## iter 10 value 20700.445870
## iter 20 value 20297.413495
## iter 30 value 20179.499867
## iter 40 value 20087.236853
## iter 50 value 20047.807320
## iter 60 value 20030.684050
## iter 70 value 20025.494688
## iter 80 value 20018.599464
## iter 90 value 20008.920657
## iter 100 value 19994.688503
## final value 19994.688503
## stopped after 100 iterations
## # weights: 8
## initial value 27291.085528
## iter 10 value 21860.071123
## iter 20 value 20734.922364
## iter 30 value 20305.923121
## iter 40 value 20201.862358
## iter 50 value 20184.717046
## iter 60 value 20183.490269
## iter 70 value 20182.666539
## iter 70 value 20182.666466
## iter 70 value 20182.666365
## final value 20182.666365
## converged
## # weights: 22
## initial value 28578.189962
## iter 10 value 20507.969706
## iter 20 value 20200.573312
## iter 30 value 20045.318776
## iter 40 value 20025.958363
## iter 50 value 20015.512556
## iter 60 value 20010.447248
## iter 70 value 20007.450291
## iter 80 value 20006.908657
## iter 90 value 20006.837712
## final value 20006.818498
## converged
## # weights: 36
## initial value 29568.695334
## iter 10 value 20186.489125
## iter 20 value 20052.673135
## iter 30 value 20036.005993
## iter 40 value 20017.034722
## iter 50 value 19991.647543

```

```

## iter 60 value 19977.313284
## iter 70 value 19974.134932
## iter 80 value 19973.556106
## iter 90 value 19971.481662
## iter 100 value 19969.042421
## final value 19969.042421
## stopped after 100 iterations
## # weights: 8
## initial value 27394.653725
## iter 10 value 21761.017847
## iter 20 value 20311.957320
## iter 30 value 20256.279602
## iter 40 value 20243.643923
## iter 50 value 20242.719628
## final value 20242.679526
## converged
## # weights: 22
## initial value 28232.840671
## iter 10 value 20399.833288
## iter 20 value 20301.551389
## iter 30 value 20178.773073
## iter 40 value 20106.224449
## iter 50 value 20071.319148
## iter 60 value 20063.873747
## iter 70 value 20052.600326
## iter 80 value 20041.130164
## iter 90 value 20040.225703
## iter 100 value 20040.102699
## final value 20040.102699
## stopped after 100 iterations
## # weights: 36
## initial value 26755.868286
## iter 10 value 20488.196791
## iter 20 value 20213.583104
## iter 30 value 20150.971938
## iter 40 value 20100.385195
## iter 50 value 20073.259782
## iter 60 value 20048.074387
## iter 70 value 20036.726175
## iter 80 value 20032.979734
## iter 90 value 20029.472485
## iter 100 value 20023.445931
## final value 20023.445931
## stopped after 100 iterations
## # weights: 8
## initial value 30758.395337
## iter 10 value 21368.400525
## iter 20 value 20331.436452
## iter 30 value 20261.873534
## iter 40 value 20253.059344
## iter 50 value 20252.341600
## final value 20252.245213
## converged
## # weights: 22

```

```

## initial value 27743.867033
## iter 10 value 20984.575038
## iter 20 value 20406.804641
## iter 30 value 20212.754940
## iter 40 value 20136.878960
## iter 50 value 20096.748767
## iter 60 value 20089.007221
## iter 70 value 20065.608823
## iter 80 value 20050.725794
## iter 90 value 20045.793333
## iter 100 value 20045.234287
## final value 20045.234287
## stopped after 100 iterations
## # weights: 36
## initial value 30654.765893
## iter 10 value 20317.006713
## iter 20 value 20145.284030
## iter 30 value 20103.054377
## iter 40 value 20082.833306
## iter 50 value 20053.000689
## iter 60 value 20031.710273
## iter 70 value 20022.302143
## iter 80 value 20018.973155
## iter 90 value 20014.448105
## iter 100 value 20012.250894
## final value 20012.250894
## stopped after 100 iterations
## # weights: 8
## initial value 26799.496453
## iter 10 value 24127.446802
## iter 20 value 20759.276338
## iter 30 value 20303.856845
## iter 40 value 20251.922162
## iter 50 value 20242.932188
## final value 20242.748199
## converged
## # weights: 22
## initial value 28436.607384
## iter 10 value 20523.096408
## iter 20 value 20301.926565
## iter 30 value 20107.020540
## iter 40 value 20054.636545
## iter 50 value 20047.887934
## iter 60 value 20044.158656
## iter 70 value 20041.482448
## iter 80 value 20039.988170
## iter 90 value 20039.360312
## iter 100 value 20039.314153
## final value 20039.314153
## stopped after 100 iterations
## # weights: 36
## initial value 29259.034977
## iter 10 value 20278.223097
## iter 20 value 20083.996320

```

```

## iter 30 value 20058.150023
## iter 40 value 20035.433609
## iter 50 value 20021.269848
## iter 60 value 20016.049555
## iter 70 value 20015.131698
## iter 80 value 20014.979065
## iter 90 value 20014.913785
## iter 100 value 20014.547775
## final value 20014.547775
## stopped after 100 iterations
## # weights: 8
## initial value 28728.267010
## iter 10 value 22578.678709
## iter 20 value 20682.057081
## iter 30 value 20333.710878
## iter 40 value 20277.340450
## iter 50 value 20264.469952
## iter 60 value 20263.743216
## iter 70 value 20263.161669
## final value 20263.139394
## converged
## # weights: 22
## initial value 26423.876103
## iter 10 value 20742.775700
## iter 20 value 20409.022711
## iter 30 value 20229.655468
## iter 40 value 20195.718463
## iter 50 value 20173.592346
## iter 60 value 20168.998209
## iter 70 value 20161.337393
## iter 80 value 20157.767869
## iter 90 value 20137.503561
## iter 100 value 20132.550851
## final value 20132.550851
## stopped after 100 iterations
## # weights: 36
## initial value 26838.286932
## iter 10 value 20177.859347
## iter 20 value 20101.331389
## iter 30 value 20089.092014
## iter 40 value 20075.109714
## iter 50 value 20061.247208
## iter 60 value 20055.357676
## iter 70 value 20052.315916
## iter 80 value 20051.618935
## iter 90 value 20050.731079
## iter 100 value 20048.978713
## final value 20048.978713
## stopped after 100 iterations
## # weights: 8
## initial value 27961.866902
## iter 10 value 23074.206347
## iter 20 value 22093.749909
## iter 30 value 21308.441384

```

```

## iter 40 value 20376.070580
## iter 50 value 20288.025260
## iter 60 value 20277.350522
## iter 70 value 20276.339863
## iter 70 value 20276.339858
## iter 70 value 20276.339858
## final value 20276.339858
## converged
## # weights: 22
## initial value 29436.882889
## iter 10 value 20758.790784
## iter 20 value 20325.369424
## iter 30 value 20259.673648
## iter 40 value 20221.256250
## iter 50 value 20150.371390
## iter 60 value 20113.630055
## iter 70 value 20098.079379
## iter 80 value 20093.608543
## iter 90 value 20091.859864
## iter 100 value 20088.827595
## final value 20088.827595
## stopped after 100 iterations
## # weights: 36
## initial value 28417.532675
## iter 10 value 21028.117121
## iter 20 value 20413.832211
## iter 30 value 20284.645720
## iter 40 value 20208.623652
## iter 50 value 20132.537347
## iter 60 value 20113.238736
## iter 70 value 20094.945354
## iter 80 value 20085.812749
## iter 90 value 20079.211312
## iter 100 value 20069.041726
## final value 20069.041726
## stopped after 100 iterations
## # weights: 8
## initial value 31172.703825
## iter 10 value 23390.607027
## iter 20 value 21235.224394
## iter 30 value 20530.020290
## iter 40 value 20292.452102
## iter 50 value 20270.385430
## iter 60 value 20263.635020
## final value 20263.594541
## converged
## # weights: 22
## initial value 27678.206007
## iter 10 value 22146.753768
## iter 20 value 21363.842619
## iter 30 value 20362.601085
## iter 40 value 20297.310530
## iter 50 value 20268.444091
## iter 60 value 20221.878841

```

```

## iter 70 value 20150.550220
## iter 80 value 20108.457002
## iter 90 value 20081.849377
## iter 100 value 20078.119056
## final value 20078.119056
## stopped after 100 iterations
## # weights: 36
## initial value 28235.757360
## iter 10 value 20294.768325
## iter 20 value 20137.976456
## iter 30 value 20095.113349
## iter 40 value 20068.819794
## iter 50 value 20055.107659
## iter 60 value 20048.431832
## iter 70 value 20045.591981
## iter 80 value 20044.647591
## iter 90 value 20043.357423
## iter 100 value 20038.291689
## final value 20038.291689
## stopped after 100 iterations
## # weights: 8
## initial value 26417.574526
## iter 10 value 21809.702085
## iter 20 value 20420.971676
## iter 30 value 20281.392885
## iter 40 value 20256.374884
## iter 50 value 20249.234152
## final value 20248.871649
## converged
## # weights: 22
## initial value 34543.455264
## iter 10 value 21993.013484
## iter 20 value 20962.354892
## iter 30 value 20354.831110
## iter 40 value 20251.706825
## iter 50 value 20234.883512
## iter 60 value 20222.114777
## iter 70 value 20198.401029
## iter 80 value 20153.977169
## iter 90 value 20136.056172
## iter 100 value 20128.621255
## final value 20128.621255
## stopped after 100 iterations
## # weights: 36
## initial value 28392.001925
## iter 10 value 21155.592459
## iter 20 value 20421.006957
## iter 30 value 20259.170983
## iter 40 value 20188.793626
## iter 50 value 20146.624953
## iter 60 value 20107.015138
## iter 70 value 20076.019654
## iter 80 value 20062.697590
## iter 90 value 20058.221915

```

```

## iter 100 value 20048.315827
## final value 20048.315827
## stopped after 100 iterations
## # weights: 8
## initial value 26964.483820
## iter 10 value 21687.606848
## iter 20 value 20672.544592
## iter 30 value 20341.787518
## iter 40 value 20267.854686
## iter 50 value 20259.078029
## final value 20259.052882
## converged
## # weights: 22
## initial value 27200.366980
## iter 10 value 20480.359266
## iter 20 value 20301.753263
## iter 30 value 20128.106457
## iter 40 value 20094.465618
## iter 50 value 20087.673461
## iter 60 value 20082.768602
## iter 70 value 20078.742362
## iter 80 value 20075.129692
## iter 90 value 20071.566508
## iter 100 value 20070.527336
## final value 20070.527336
## stopped after 100 iterations
## # weights: 36
## initial value 26851.054653
## iter 10 value 20243.155717
## iter 20 value 20128.105465
## iter 30 value 20093.649094
## iter 40 value 20066.261179
## iter 50 value 20046.389658
## iter 60 value 20038.163040
## iter 70 value 20035.242229
## iter 80 value 20034.815981
## iter 90 value 20034.022969
## iter 100 value 20033.376720
## final value 20033.376720
## stopped after 100 iterations
## # weights: 8
## initial value 29000.770221
## iter 10 value 20612.791439
## iter 20 value 20299.245532
## iter 30 value 20256.861059
## iter 40 value 20250.219237
## iter 50 value 20248.683785
## final value 20248.619786
## converged
## # weights: 22
## initial value 28967.406723
## iter 10 value 20504.782459
## iter 20 value 20286.757548
## iter 30 value 20186.310331

```

```

## iter 40 value 20135.619936
## iter 50 value 20108.050960
## iter 60 value 20097.344965
## iter 70 value 20078.557222
## iter 80 value 20074.306124
## iter 90 value 20072.974830
## iter 100 value 20072.852174
## final value 20072.852174
## stopped after 100 iterations
## # weights: 36
## initial value 27260.178636
## iter 10 value 20399.362113
## iter 20 value 20185.258552
## iter 30 value 20140.727802
## iter 40 value 20110.054317
## iter 50 value 20095.787769
## iter 60 value 20069.422029
## iter 70 value 20058.624788
## iter 80 value 20052.152169
## iter 90 value 20047.358189
## iter 100 value 20042.351225
## final value 20042.351225
## stopped after 100 iterations
## # weights: 8
## initial value 27768.247847
## iter 10 value 20481.578592
## iter 20 value 20197.048021
## iter 30 value 20180.175991
## iter 40 value 20178.683957
## final value 20178.667094
## converged
## # weights: 22
## initial value 26481.298699
## iter 10 value 21045.872157
## iter 20 value 20526.573100
## iter 30 value 20246.022348
## iter 40 value 20143.940796
## iter 50 value 20108.480705
## iter 60 value 20085.647820
## iter 70 value 20008.718201
## iter 80 value 20001.815950
## iter 90 value 20000.773770
## iter 100 value 20000.706448
## final value 20000.706448
## stopped after 100 iterations
## # weights: 36
## initial value 31520.750516
## iter 10 value 20292.736041
## iter 20 value 20067.481801
## iter 30 value 20018.950754
## iter 40 value 19984.779088
## iter 50 value 19958.421041
## iter 60 value 19953.126761
## iter 70 value 19951.890368

```



```

## iter 80 value 19951.383593
## iter 90 value 19950.739109
## iter 100 value 19950.342561
## final value 19950.342561
## stopped after 100 iterations
## # weights: 8
## initial value 28288.768356
## iter 10 value 21497.100313
## iter 20 value 20570.040488
## iter 30 value 20232.328811
## iter 40 value 20197.167477
## iter 50 value 20190.473690
## final value 20190.473428
## converged
## # weights: 22
## initial value 28198.760369
## iter 10 value 21239.549769
## iter 20 value 20911.372536
## iter 30 value 20217.577771
## iter 40 value 20044.881982
## iter 50 value 20016.006910
## iter 60 value 20006.452217
## iter 70 value 19995.729392
## iter 80 value 19992.617283
## iter 90 value 19992.190565
## final value 19992.180839
## converged
## # weights: 36
## initial value 26960.023341
## iter 10 value 21030.877584
## iter 20 value 20479.285342
## iter 30 value 20324.006976
## iter 40 value 20189.442506
## iter 50 value 20065.547737
## iter 60 value 20040.515303
## iter 70 value 20010.568595
## iter 80 value 19992.952141
## iter 90 value 19982.331133
## iter 100 value 19974.365465
## final value 19974.365465
## stopped after 100 iterations
## # weights: 8
## initial value 27135.907439
## iter 10 value 22643.305938
## iter 20 value 21561.796336
## iter 30 value 21238.666530
## iter 40 value 20401.085043
## iter 50 value 20203.230040
## iter 60 value 20184.650192
## iter 70 value 20178.572394
## final value 20178.542531
## converged
## # weights: 22
## initial value 26476.658834

```

```

## iter 10 value 22172.890074
## iter 20 value 21323.448278
## iter 30 value 20364.833195
## iter 40 value 20051.178872
## iter 50 value 20022.072650
## iter 60 value 20009.699106
## iter 70 value 19991.989419
## iter 80 value 19986.910416
## iter 90 value 19986.425504
## final value 19986.419835
## converged
## # weights: 36
## initial value 29175.567742
## iter 10 value 20209.096423
## iter 20 value 20039.281980
## iter 30 value 20015.427048
## iter 40 value 19984.719593
## iter 50 value 19974.605185
## iter 60 value 19968.525536
## iter 70 value 19964.452072
## iter 80 value 19962.976152
## iter 90 value 19960.281971
## iter 100 value 19958.482026
## final value 19958.482026
## stopped after 100 iterations
## # weights: 36
## initial value 33172.352904
## iter 10 value 25218.647485
## iter 20 value 25114.728115
## iter 30 value 25088.228096
## iter 40 value 25059.191615
## iter 50 value 25035.173589
## iter 60 value 25029.659170
## iter 70 value 25024.883167
## iter 80 value 25020.378806
## iter 90 value 25017.091219
## iter 100 value 25010.702772
## final value 25010.702772
## stopped after 100 iterations

```

```
summary(model)
```

```

## a 5-5-1 network with 36 weights
## options were - entropy fitting
## b->h1 i1->h1 i2->h1 i3->h1 i4->h1 i5->h1
## -1.91 -0.46 -0.43 -1.35 -0.86 0.29
## b->h2 i1->h2 i2->h2 i3->h2 i4->h2 i5->h2
## -1.24 -1.06 -1.20 0.31 0.14 0.03
## b->h3 i1->h3 i2->h3 i3->h3 i4->h3 i5->h3
## 1.29 -0.16 0.21 0.05 0.15 0.56
## b->h4 i1->h4 i2->h4 i3->h4 i4->h4 i5->h4
## -6.51 -0.13 -0.96 -5.95 -2.18 0.69
## b->h5 i1->h5 i2->h5 i3->h5 i4->h5 i5->h5
## -1.38 -4.28 1.07 0.18 -1.41 -2.77
## b->o h1->o h2->o h3->o h4->o h5->o

```

```

## -4.30 -4.68 -1.92  7.33  2.05 -0.34
test_set$pred <- predict(model, test_set)
test_set$factor_pred <- as.factor(test_set$pred)
test_set$factor_truth <- as.factor(test_set$Diabetes_binary)
precision <- posPredValue(test_set$factor_truth, test_set$factor_pred)
recall <- sensitivity(test_set$factor_truth, test_set$factor_pred)
cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table
test_set$pred <- predict(model, test_set, probability=TRUE)
#-----
print(paste("Precision: ", precision))

## [1] "Precision:  0.695425991362387"

print(paste("Accuracy: ", accuracy))

## [1] "Accuracy:  0.743254332190954"


print(paste("Recall: ", recall))

## [1] "Recall:  0.761746048811956"

acc_nn <- accuracy
#-----
Reference <- factor(c(0, 1, 0, 1))
Prediction <- factor(c(0, 0, 1, 1))
Y <- array(cmat)
df <- data.frame(Prediction, Reference, Y)

library(ggplot2)
ggplot(data = df, mapping = aes(x = Reference, y = Prediction)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%1.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow")

```



```
model <- train(Diabetes_binary ~ Age+GenHlth+HighChol+HighBP+BMI, method='ctree', data = training_set)
summary(model)
```

```
##      Length      Class      Mode
##      1 BinaryTree      S4
```

```
test_set$pred <- predict(model, test_set)
test_set$factor_pred <- as.factor(test_set$pred)
test_set$factor_truth <- as.factor(test_set$Diabetes_binary)
precision <- posPredValue(test_set$factor_truth, test_set$factor_pred)
recall <- sensitivity(test_set$factor_truth, test_set$factor_pred)
cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table
test_set$pred <- predict(model, test_set, probability=TRUE)
# -----
print(paste("Precision: ", precision))
```

```
## [1] "Precision: 0.677856301531213"
```

```
print(paste("Accuracy: ", accuracy))
```

```
## [1] "Accuracy: 0.73326253801226"
```

```
print(paste("Recall: ", recall))
```

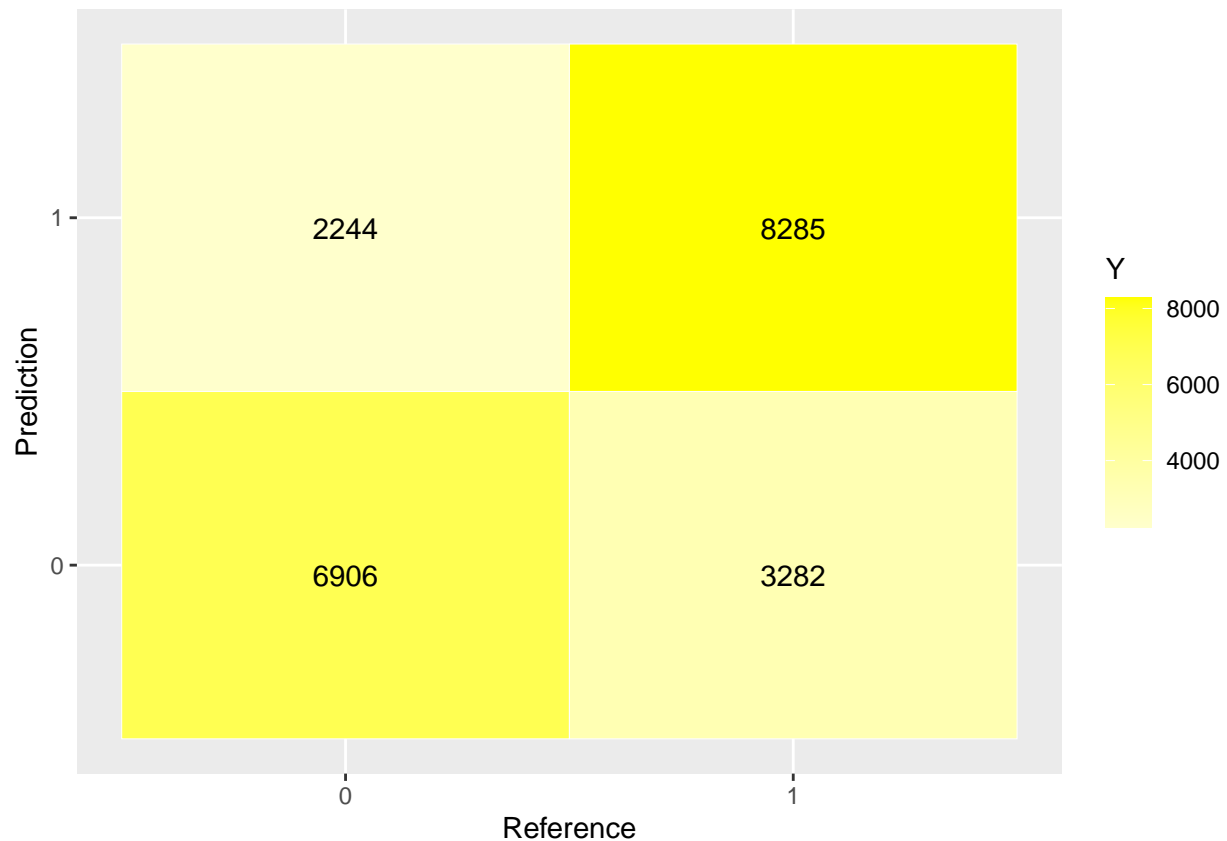
```
## [1] "Recall: 0.754754098360656"
```

```

acc_tree <- accuracy
# -----
Reference <- factor(c(0, 1, 0, 1))
Prediction <- factor(c(0, 0, 1, 1))
Y <- array(cmat)
df <- data.frame(Prediction, Reference, Y)

library(ggplot2)
ggplot(data = df, mapping = aes(x = Reference, y = Prediction)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%1.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow")

```



```

model <- train(Diabetes_binary ~ Age+GenHlth+HighChol+HighBP+BMI, method='svmLinear', data = training_s
summary(model)

```

```

## Length Class Mode
##      1  ksvm   S4

test_set$pred <- predict(model, test_set)
test_set$factor_pred <- as.factor(test_set$pred)
test_set$factor_truth <- as.factor(test_set$Diabetes_binary)
precision <- posPredValue(test_set$factor_truth, test_set$factor_pred)
recall <- sensitivity(test_set$factor_truth, test_set$factor_pred)
cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table

```

```

test_set$pred <- predict(model, test_set,probability=TRUE)
#-----
print(paste("Precision: ",precision))

## [1] "Precision:  0.686788378484492"

print(paste("Accuracy: ",accuracy))

## [1] "Accuracy:  0.737799874499204"


print(paste("Recall: ",recall))

## [1] "Recall:  0.757415024897164"

acc_svm <- accuracy
#-----
Reference <- factor(c(0, 1, 0, 1))
Prediction <- factor(c(0, 0, 1, 1))
Y <- array(cmat)
df <- data.frame(Prediction, Reference, Y)

library(ggplot2)
ggplot(data = df, mapping = aes(x = Reference, y = Prediction)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%1.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow")

```



```
model <- train(Diabetes_binary ~ Age+GenHlth+HighChol+HighBP+BMI, method='LogitBoost', data = training_data)
summary(model)
```

```
##           Length Class      Mode
## Stump      33    -none-    numeric
## lablist     2     factor    numeric
## xNames      5    -none-    character
## problemType 1    -none-    character
## tuneValue   1    data.frame list
## obsLevels   2    -none-    character
## param       0    -none-    list
```

```
test_set$pred <- predict(model, test_set)
test_set$factor_pred <- as.factor(test_set$pred)
test_set$factor_truth <- as.factor(test_set$Diabetes_binary)
precision <- posPredValue(test_set$factor_truth, test_set$factor_pred)
recall <- sensitivity(test_set$factor_truth, test_set$factor_pred)
cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table
test_set$pred <- predict(model, test_set, probability=TRUE)
#-----
print(paste("Precision: ", precision))
```

```
## [1] "Precision: 0.742343934040047"
```

```
print(paste("Accuracy: ", accuracy))
```

```
## [1] "Accuracy: 0.712603176135541"
```

```
print(paste("Recall: ", recall))
```

```
## [1] "Recall: 0.694362835108336"
```

```
acc_lr <- accuracy
```

```
#-----
```

```
Reference <- factor(c(0, 1, 0, 1))
```


```
Prediction <- factor(c(0, 0, 1, 1))
```

```
Y <- array(cmat)
```

```
df <- data.frame(Prediction, Reference, Y)
```

```
library(ggplot2)
```

```
ggplot(data = df, mapping = aes(x = Reference, y = Prediction)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow")
```



```
model <- train(Diabetes_binary ~ Age+GenHlth+HighChol+HighBP+BMI, method='knn', data = training_set, me
summary(model)
```

```
##           Length Class      Mode
## learn      2      -none-    list
## k          1      -none-    numeric
## theDots    0      -none-    list
## xNames     5      -none-    character
## problemType 1      -none-    character
## tuneValue  1      data.frame list
## obsLevels  2      -none-    character
## param      0      -none-    list
```

```
test_set$pred <- predict(model, test_set)
test_set$factor_pred <- as.factor(test_set$pred)
test_set$factor_truth <- as.factor(test_set$Diabetes_binary)
precision <- posPredValue(test_set$factor_truth, test_set$factor_pred)
recall <- sensitivity(test_set$factor_truth, test_set$factor_pred)
cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table
test_set$pred <- predict(model, test_set, probability=TRUE)
# -----
print(paste("Precision: ", precision))
```

```
## [1] "Precision: 0.691794267765999"
```



```

print(paste("Accuracy: ",accuracy))

## [1] "Accuracy:  0.730511174397838"


print(paste("Recall: ",recall))

## [1] "Recall:  0.742598250974608"

acc_knn <- accuracy
# -----
Reference <- factor(c(0, 1, 0, 1))
Prediction  <- factor(c(0, 0, 1, 1))
Y <- array(cmat)
df <- data.frame(Prediction, Reference, Y)

library(ggplot2)
ggplot(data = df, mapping = aes(x = Reference, y = Prediction)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%1.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow")

```



```

print(paste('Accuracy in QDA:',acc_qda*100,'%'))

## [1] "Accuracy in QDA: 73.3793502920307 %"

print(paste('Accuracy in LDA:',acc_lda*100,'%'))

## [1] "Accuracy in LDA: 73.6496596997635 %"

```

```

print(paste('Accuracy in Randomforrest:',acc_rf*100,'%'))

## [1] "Accuracy in Randomforrest: 73.9441038760438 %"
print(paste('Accuracy in Neural Network:',acc_nn*100,'%'))

## [1] "Accuracy in Neural Network: 74.3254332190954 %"
print(paste('Accuracy in CTree:',acc_tree*100,'%'))

## [1] "Accuracy in CTree: 73.326253801226 %"
print(paste('Accuracy in Linear SVM:',acc_svm*100,'%'))

## [1] "Accuracy in Linear SVM: 73.7799874499204 %"
print(paste('Accuracy in Logreg:',acc_lr*100,'%'))

## [1] "Accuracy in Logreg: 71.2603176135541 %"
print(paste('Accuracy in knn:',acc_knn*100,'%'))

## [1] "Accuracy in knn: 73.0511174397838 %"

```

PART 5

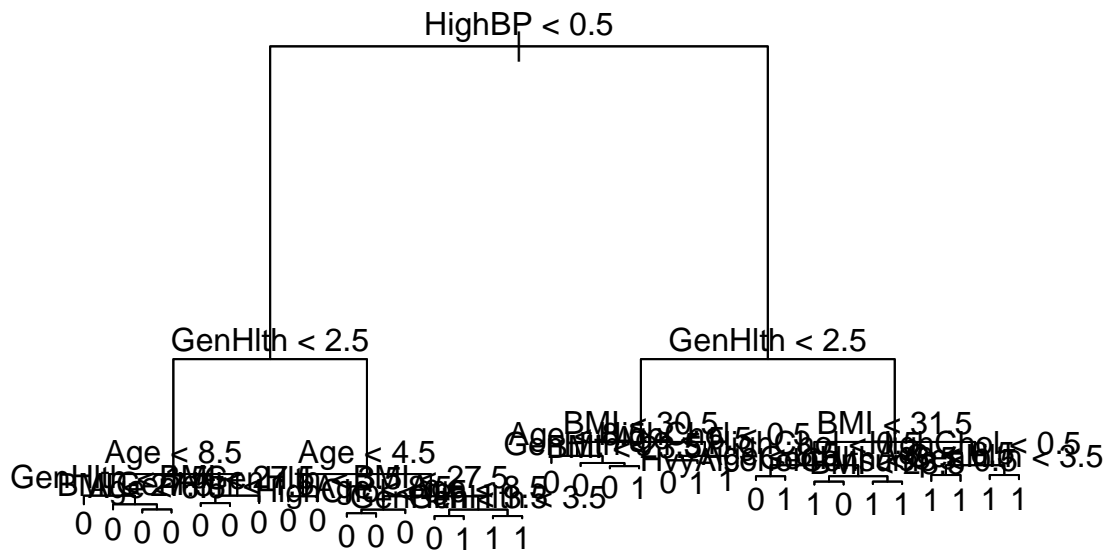
Yes, we can do this thing. For this target, we need to use a simple and appropriate inferentially. As we want to ask some questions as less as possible from a person, I prefer to use the Tree model. In the last part, we understood Tree can do nice predictions too. so use a big tree and use cross-validation to get how many nodes and terminals we need for a good model. So we choose 8 and get accuracy, a plot of the Tree, and a confusion matrix for this model. By this tree, we can by asking just some questions, we can predict whether one has diabetes or not, with acceptable accuracy.

```

tree_model <- tree(Diabetes_binary ~ .-output , training_set , mindev=7e-4, minsize=8)
summary(tree_model)

##
## Classification tree:
## tree(formula = Diabetes_binary ~ . - output, data = training_set,
##      mindev = 7e-04, minsize = 8)
## Variables actually used in tree construction:
## [1] "HighBP"          "GenHlth"          "Age"
## [4] "BMI"             "HighChol"         "HvyAlcoholConsump"
## Number of terminal nodes: 33
## Residual mean deviance: 1.064 = 51400 / 48310
## Misclassification error rate: 0.2685 = 12978 / 48340
plot(tree_model)
text(tree_model, pretty = 0)

```




```

test_set$pred <- predict(tree_model, test_set, type="class")
ErrorPrune<-mean(test_set$pred!=test_set$Diabetes_binary)
cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table

real <- factor(c(0, 0, 1, 1))
pred <- factor(c(0, 1, 0, 1))
Y <- array(cmat)
df <- data.frame(pred, real, Y)

ggplot(data = df, mapping = aes(x = real, y = pred)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%1.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow") +
  theme_bw()

```



```
print(paste('Error rate :',ErrorPrune))
```

```
## [1] "Error rate : 0.269488825602162"
```

```
cv.tree<-cv.tree(tree_model, FUN=prune.tree)
```

```
cv.tree
```

```
## $size
```

```
## [1] 33 32 31 30 29 28 27 26 25 24 22 21 20 19 18 17 16 15 14 13 12 11 10 9 8
```

```
## [26] 7 6 5 4 3 2 1
```

```
##
```

```
## $dev
```

```
## [1] 55975.35 55975.35 55975.35 55975.35 55975.35 55975.35 55975.35 55975.35
```

```
## [9] 55975.35 55975.35 55975.35 55975.35 55975.35 55975.35 55975.35 55975.35
```

```
## [17] 55975.35 55975.35 55975.35 55975.35 55975.35 55975.35 55975.35 55975.35
```

```
## [25] 55975.35 55975.35 55975.35 55975.35 55975.35 57633.17 60364.83 67002.35
```

```
##
```

```
## $k
```

```
## [1] -Inf 48.01600 48.50432 53.19121 62.48372 65.39000
```

```
## [7] 65.59553 67.79974 68.82130 69.10782 77.24479 80.21836
```

```
## [13] 89.66217 95.64933 110.37263 111.01990 115.88648 129.18833
```

```
## [19] 138.87016 154.00557 161.24952 186.67753 212.97755 238.36911
```

```
## [25] 289.57154 300.56292 466.27489 479.05378 502.23111 1789.48871
```


```
## [31] 2483.36772 6759.14441
```

```
##
```


```
## $method
```

```
## [1] "deviance"
```

```
##
## attr("class")
## [1] "prune"          "tree.sequence"
plot(cv.tree)
```



```
prune.tree<-prune.misclass(tree_model,best = 8)
plot(prune.tree);text(prune.tree)
```



```

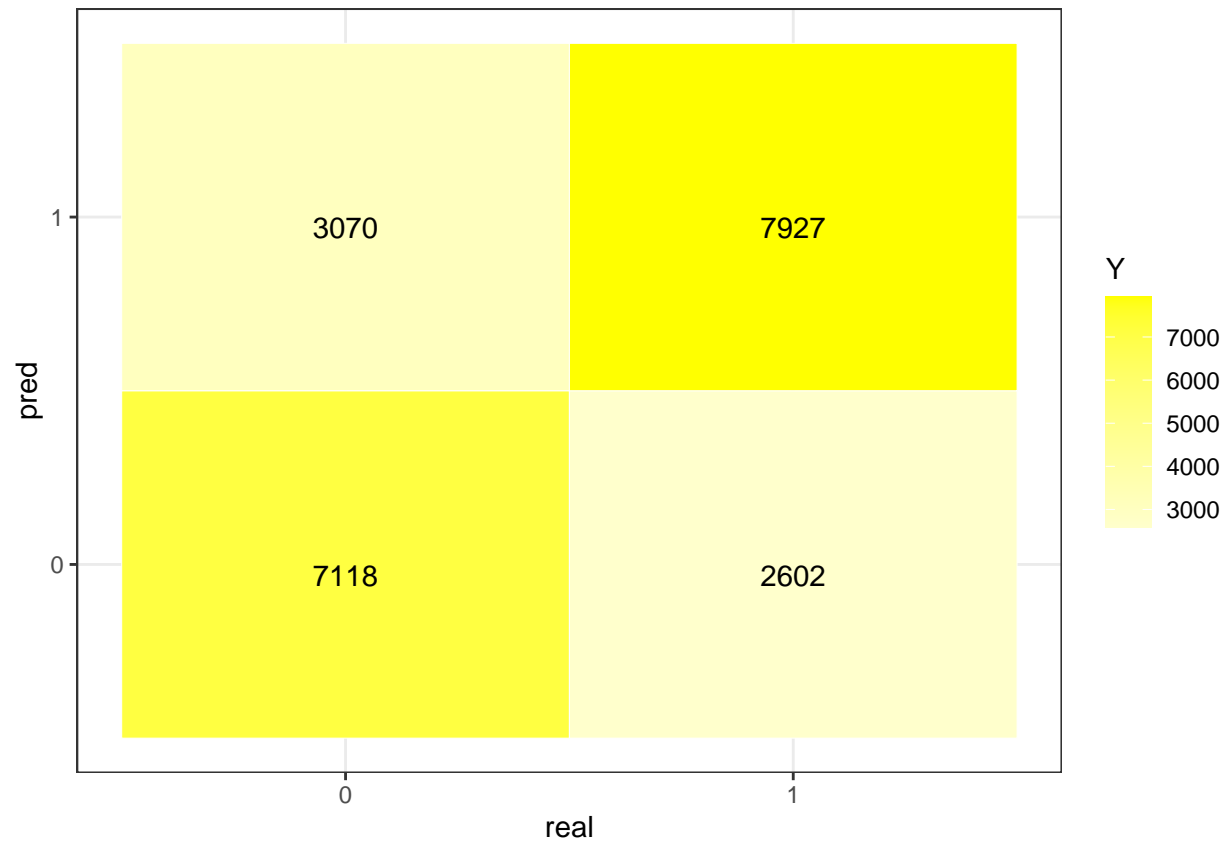
test_set$pred <- predict(prune.tree, test_set, type="class")
ErrorPrune <- mean(test_set$pred != test_set$Diabetes_binary)

cm <- confusionMatrix(test_set$pred, test_set$Diabetes_binary)
accuracy <- cm$overall[1]
cmat <- cm$table

real <- factor(c(0, 0, 1, 1))
pred <- factor(c(0, 1, 0, 1))
Y <- array(cmat)
df <- data.frame(pred, real, Y)

library(ggplot2)
ggplot(data = df, mapping = aes(x = real, y = pred)) +
  geom_tile(aes(fill = Y), colour = "white") +
  geom_text(aes(label = sprintf("%.0f", Y)), vjust = 1) +
  scale_fill_gradient(low = "#ffffcc", high = "yellow") +
  theme_bw()

```



```
print(paste('New Error rate :',ErrorPrune))
```

```
## [1] "New Error rate : 0.27378481440363"
```

```
print(paste('New accuracy :',((1-ErrorPrune)*100), '%'))
```

```
## [1] "New accuracy : 72.621518559637 %"
```